

1) Markov Process

Tuple (S, P, ss') = (set of states, state transition proba. matrix)
 $P_{ss'} = P[S_{t+1} = s' | S_t = s]$

Generates chain of Markov states governed by prob. transitions
 A state s_t is **Markov** if and only if $P[S_{t+1} | S_t] = P[S_{t+1} | s_1, \dots, s_t]$

Transient states (round) vs terminal states (square)

Stationary/homogeneous chain: $P_{ss'}$ only depends on s, s' (not t)

MRP: Markov chain which emits rewards (S, P, R, γ)

$\mathcal{R}_s = \mathbb{E}[r_{t+1} | S_t = s]$ Expected immediate reward collected upon departing s , at $t+1$

$\gamma \in [0, 1]$ Discount factor (see advt.)

$R_t = r_{t+1} + \gamma r_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$ Return = total discounted reward from t

The factor $\gamma \in [0, 1]$ is how we **discount** the present value of future rewards.
 The value of receiving reward r after $k+1$ time-steps is $\gamma^k r$.
 The discount values immediate reward higher than delayed reward:
 • γ close to 0 leads to "myopic" (short-sighted) evaluation.
 • γ close to 1 leads to "far-sighted" evaluation.

State-value fn = expected return R starting from state s at time t
 $v(s) = \mathbb{E}[R_t | S_t = s]$

Bellman Eq for MRPs (start from above)

$v(s) = \mathbb{E}[r_{t+1} + \gamma v(S_{t+1}) | S_t = s]$
 $= \mathcal{R}_s + \gamma \sum_{s' \in S} P_{ss'} v(s')$ If $|S|=n$, n eqs

$\mathbf{v} = \mathcal{R} + \gamma \mathbf{Pv}$ If $|S|=n$, n -dim vector \mathbf{v}

Bellman Eq = linear, self-consistent \Rightarrow can solve \mathbf{v} directly if small MRP (matrix inv.). **Iterative Methods** for larger MRPs: DP/MC/TD

Probabilistic/stochastic or deterministic policy

A policy $\pi_t(a, s) = P[A_t = a | S_t = s]$ is the conditional probability distribution to execute an action $a \in \mathcal{A}$ given that one is in state $s \in S$ at time t .

Deterministic: $\pi_t(s) = a$ vs stochastic: $\mathbf{a} \sim p_{\pi}(\mathbf{a} | \mathbf{s})$
 $p_{\pi}(\mathbf{a} | \mathbf{s}) = \pi(\mathbf{a} | \mathbf{s}) \equiv \pi(a, s)$

4) MC (for V estimation)

Learn from complete episodes of sample traces (no bootstrapping). $V(s)$ = mean over empirical returns observed after visits to s (not expected R ; will converge to it) = MC Policy Evaluation by sampling values for V . Incremental updates MC.

First-visit vs Every-visit MC: record return of episode E from 1st occurrence vs every occurrence of s in E , for all s .

Vanilla MC: update pi *once* after X episodes. **Batch MC**: update pi every *batch_size* episodes. **Online MC**: update pi *every* episode

We can now update value functions without having to store sample traces:

- Update $V(s)$ incrementally after step $s_t, a_t, r_{t+1}, S_{t+1}$
- For each state s_t with return R_t (up to this point) and $N(s)$ the visit counter to this state:

$$N(s_t) \leftarrow N(s_t) + 1$$

$$V(s_t) \leftarrow V(s_t) + \frac{1}{N(s_t)} (R_t - V(s_t))$$

Moreover, if the world is **non-stationary**, it can be useful to track a **running mean**, i.e. by gradually forgetting old episodes.

MC update at t $V(s_t) \leftarrow V(s_t) + \alpha (R_t - V(s_t))$ **Alpha = $1/r$**

The parameter α controls the rate of forgetting old episodes

Advt:

- model-free (no MDP knowledge needed: R/P) since R sampled
- fights curseOfDim through sampling (sample backups)
- cost of backups = constant (not exp.) + indept of $N=|S|$
- zero bias
- good convergence properties (even w/ FA)
- not v sensitive to initial value
- v simple to understand and use

Disv:

- only episodic MDPs w/ terminal states
- high variance: sampled R depends on many random A, P, R

Usually more effective in non-Markov env (no Markov assumed); TD exploits Markov property so TD more efficient in Markov env.

Model-free prediction (policy eval): estimate v fn of unknown MDP

Model-free control (policy improv): optimise v fn \Rightarrow see 6

TD control: on-policy (**SARSA**) vs off-policy (**Q-learning**)

Apply TD to $Q(S, A)$ + use e-greedy policy improv. + update every t

SARSA: Use GPI like in DP/MC, but eval./pred. w/ TD on Q (conv.)
 $Q(S, A) \leftarrow Q(S, A) + \alpha (r + \gamma Q(S', A') - Q(S, A))$


Q-learning (w/ assumpt. of coverage): improve pi (greedy wrt $Q(s, a)$) and pi' (e-greedy wrt $Q(s, a)$). Q-learning update (R = immediate r):

$$Q(S, A) \leftarrow Q(S, A) + \alpha (R + \gamma \max_{a'} Q(S', a') - Q(S, A))$$

No explicit policies in algo: pi implicit in greedy term; pi' = e-greedy version of pi. Both updated every t bc Q updated every t .

TD(0) here: immediate r + 1-step-look-ahead with $Q(s', a')$

2) MDP




V^{π} has unique sol

Value fn (self-consistent, linear)

$$V^{\pi}(s) = \mathbb{E}_{\pi}[R_t | S_t = s] = \mathbb{E}[\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} | S_t = s]$$

Backup Diagram: transfer state value info from all s' to s (update/backup op)



To compute $v(s)$: average over all possible traces and their reward

Another version of Bellman Eq (BEq)

$$V^{\pi}(s) = \sum_{a \in \mathcal{A}} \pi(s, a) \sum_{s' \in S} P_{ss'}^a (\mathcal{R}_{ss'}^a + \gamma V^{\pi}(s'))$$

Value fns satisfy a set of recursive consistency eqs.

State-Action value fn (cost-to-go)


$$Q^{\pi}(s, a) = \mathbb{E}[R_t | S_t = s, A_t = a] = \mathbb{E}[\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} | S_t = s, A_t = a]$$

Ordering of policies: $\pi \geq \pi'$ iff $V^{\pi}(s) \geq V^{\pi'}(s)$ for all $s \in S$

Optimal V^* : $V^*(s) = \max_{\pi} V^{\pi}(s), \forall s \in S$ **Optimal π^*** : Therefore, the policy π^* that maximises the value function is the **optimal policy**. There is always at least one optimal policy. There may be more than one optimal policy.

Optimal Q^* : $Q^*(s, a) = \max_{\pi} Q^{\pi}(s, a), \forall s \in S, a \in \mathcal{A} = \mathbb{E}[r_{t+1} + \gamma V^*(S_{t+1}) | S_t = s, A_t = a]$

Backup diagram for Q^{π} pi(s,a)



$$V^*(s) = \max_a \sum_{s'} P[s' | s, a] (r(s, a, s') + \gamma V^*(s'))$$

$$= \max_a \sum_{s'} P_{ss'}^a (\mathcal{R}_{ss'}^a + \gamma V^*(s'))$$

BOEs: **direct sol** to find π^* but hard + relies on assumptions: know env dynamics; comp resources; Markov property.

BOEq Convergence Theorem \Rightarrow

relationship btw Q and V

$$V^{\pi}(s) = \sum_{a \in \mathcal{A}} \pi(s, a) Q^{\pi}(s, a)$$

if $|S|=n$, there are n V^* eqs; unique sol V^* indpt of pi.

$$Q^*(s, a) = \mathbb{E}[r_{t+1} + \gamma \max_{a'} Q^*(S_{t+1}, a') | S_t = s, A_t = a]$$

$$= \sum_{s'} P_{ss'}^a (\mathcal{R}_{ss'}^a + \gamma \max_{a'} Q^*(s', a'))$$


For an MDP with a finite state and action space

- The Bellman (Optimality) equations have a unique solution.
- The values produced by value iteration converge to the solution of the Bellman equations.

3) DP (solve/find π^* known MDPs) Assume: finite MDP + know perfect model of env (P & R) + prblm w/ optimal substructure and overlapping subprblms. Simplify prblm by breaking it down into simpler subprblms recursively (Principle of Optimality; BEq as relation btw value of larger prblm & values of subprblms).

Advt:

- can do synchr (all states backed up in V): 2 copies of V or asynchr (1 copy of V , in-place; sig. reduced computation + still conv. guaranteed) updates
- bootstrapping: efficient use of data (thx optimal substructure)
- Disv:**
- model-based
- curse of dimensionality ($>$ mil. states)
- even 1 backup too costly



Policy Improvement Theorem

Let π and π' be any two deterministic policies such that $\forall s \in S$:

$$Q^{\pi'}(s, \pi'(s)) \geq V^{\pi}(s)$$

Then π' must be as good or better than π :

$$V^{\pi'}(s) \geq V^{\pi}(s), \forall s \in S$$

Policy Iteration (PI): conv. guaranteed; eval \leftrightarrow imprv

Once a policy, π , has been improved using V^{π} to yield a better policy π' , we can compute $V^{\pi'}$ and improve it again to π'' , to yield an even better $V^{\pi''}$. We can thus obtain a sequence of monotonically improving policies and value functions. This way of finding an optimal policy is called **policy iteration**.

Principle of Optimality

A policy $\pi(a|s)$ achieves the optimal value from state s , $V^{\pi}(s) = V^*(s)$, if and only if

- For any state s' reachable from s , i.e. $\exists a: p(s', s, a) > 0$
- π achieves the optimal value from state s' , $V^{\pi}(s') = V^*(s')$

Value Iteration (VI): PI w/ 1 iteration only of pi eval

BOEq = update rule applied iteratively (1-step look ahead)

$$V^*(s) \leftarrow \max_a \sum_{s'} P_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V^*(s')]$$

Output a deterministic policy, π , such that $\pi(s) = \arg \max_a \sum_{s'} P_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V^*(s')]$

Unlike PI, no explicit policy pi:

Incremental update after each t

$$V(s_t) \leftarrow V(s_t) + \alpha (r_{t+1} + \gamma V(s_{t+1}) - V(s_t))$$

We update the value of $V(s_t)$ towards the **estimated** return

$r_{t+1} + \gamma V(s_{t+1})$ Note, how we are combining a measurement r_{t+1} with an estimate $V(s_{t+1})$ to produce a better estimate $V(s_t)$.

$r_{t+1} + \gamma V(s_{t+1}) - V(s_t)$ is the **Temporal Difference Error**

$r_{t+1} + \gamma V(s_{t+1})$ is the **Temporal Difference Target**

Disv:

- some bias: TD target is biased estimate of $V^{\pi}(s)$ as it relies on estimate of state s_{t+1} (bootstrapping)
- convergence not guaranteed w/ FA
- more sensitive to initial value than MC

Bootstrapping = update involves an estimate

Sampling = update does not involve an expected value

6) Model-Free Control: V/Q learned by MC/TD & follow GPI Algo w/ approximate V/Q and pi converging to π^*

MC Policy Improvement: make pi (e-)greedy wrt current Q fn (no model needed to build greedy pi since Q , not V)

For any action-value function $Q(s, a)$, the corresponding greedy policy is the one that, for each $s \in S$ deterministically chooses an action with maximal action-value: $\pi(s) = \arg \max_a Q(s, a)$

For the approx. $V/Q \rightarrow$ true V/Q and the Policy Improv. Theorem w/ MC to work, assume: inf. num of episodes + exploring starts (random s_0). Then, MC can find π^* given only sample episodes + no knowledge of env dynamics

\Rightarrow **avoid exploring starts assumption**: ensure agent continues to select them \Rightarrow on/off policy methods

On-policy (learn pi from exp sampled from pi) vs **Off-policy** (learn **target pi** from exp sampled from **behavior pi'**)

on-policy method \rightarrow soft policies \rightarrow epsilon-greedy policies) e-greedy policy with $\epsilon \in [0, 1]$.

Soft policies have in general $\pi(a, s) > 0 \forall s \in S, \forall a \in \mathcal{A}$. I.e. we have a finite probability to **explore** all actions.

e.g., on-policy first-visit MC control (for e-soft policies), batch or iterative learning for control

$$\pi(s, a) = \begin{cases} 1 - \epsilon + \frac{\epsilon}{|A(s)|}, & \text{if } a^* = \arg \max_a Q(s, a) \\ \frac{\epsilon}{|A(s)|}, & \text{if } a \neq a^* \end{cases}$$

7) Function Approx. (FA)

Prblm (tabular RL): too many states/actions (CurseOfDim) to be storing values of $V(s)$ and $Q(s,a)$ as lookup tables.
SOL: estimate V/Q fn w/ FA. Generalise from seen states to unseen states. Update param w using MC/TD learning.

$$V^\pi(s) \approx \hat{V}(s, \mathbf{w}) \quad \text{FA: ANN, decision tree, nearest-neighbour...}$$
$$Q^\pi(s, a) \approx \hat{Q}(s, a, \mathbf{w})$$

Assumpt: differentiable fn, non-stationary & non-iid data
GD goal: find \mathbf{w} minimizing MSE btw approx. $\hat{V}(s, \mathbf{w})$ and true $V^\pi(s)$ fn

$$J(\mathbf{w}) = \mathbb{E} \left[(V^\pi(s) - \hat{V}(s, \mathbf{w}))^2 \right]$$
$$\Delta \mathbf{w} = -\frac{1}{2} \alpha \nabla_{\mathbf{w}} J(\mathbf{w})$$


SGD samples the gradient & the average update = full update
 $\Delta \mathbf{w} = \alpha (V^\pi(s) - \hat{V}(s, \mathbf{w})) \nabla_{\mathbf{w}} \hat{V}(s, \mathbf{w})$

States represented as hand-engineered feature vectors $\mathbf{x}(s)$

1) **Coarse coding:** represent s w/ overlapping binary features (if s in a circle, 1, else 0).

e.g. **tile coding** (suited for computer, efficient on-line learning)

2) **Radial Basis Fns (RBF):** generalise coarse coding to continuous-valued features $\{[0,1]\}$.



MC/TD for evaluation (V/Q FA)

Pseudocode for MC learning \mathbf{w}

Loop forever (for each episode):
Generate an episode $S_0, A_0, R_1, S_1, A_1, \dots, R_T, S_T$ using π
Loop for each step of episode, $t = 0, 1, \dots, T-1$:
 $\mathbf{w} \leftarrow \mathbf{w} + \alpha [G_t - \hat{V}(S_t, \mathbf{w})] \nabla \hat{V}(S_t, \mathbf{w})$

Return R_t is an unbiased, noisy sample of true value $V^\pi(S_t)$

We apply supervised learning on "training data" of state return trace:

$$(s_1, r_1), (s_2, r_2), \dots, (s_T, r_T) \quad (7)$$

For example, using linear Monte-Carlo policy evaluation

$$\Delta \mathbf{w} = \alpha (R_t - \hat{V}(s, \mathbf{w})) \nabla_{\mathbf{w}} \hat{V}(s, \mathbf{w}) \quad (8)$$

$$= \alpha (R_t - \hat{V}(s, \mathbf{w})) \mathbf{x}(s_t) \quad (9)$$

Monte-Carlo evaluation converges to a local optimum, even when using non-linear value function approximation (provable)

MC/TD FA = linear V fn approx.:

$$\hat{V}(s, \mathbf{w}) = \mathbf{x}(s)^T \mathbf{w} = \sum_{j=1}^n x_j(s) w_j$$

Update rule = $\text{lr} \times \text{pred error} \times \text{feature value } \mathbf{x}$

$$\phi_s(i) = \exp \left(-\frac{|s - c_i|^2}{2\sigma_i^2} \right)$$

where each $\phi_s(i)$ is the i -th basis function that maps a continuous state space variable s . Each RBF is centered at location c_i and has width σ_i .

The TD-target $R_{t+1} + \gamma \hat{V}(s_{t+1}, \mathbf{w})$ is a biased (single) sample of the true value $V^\pi(s_t)$

We still perform supervised learning on "digested" training data:

$$(s_1, r_2 + \gamma \hat{V}(s_2, \mathbf{w})), (s_2, r_3 + \gamma \hat{V}(s_3, \mathbf{w})), \dots, (s_T, r_T) \quad (10)$$


For example, using linear TD

$$\Delta \mathbf{w} = \alpha (r + \gamma \hat{V}(s', \mathbf{w}) - \hat{V}(s, \mathbf{w})) \nabla_{\mathbf{w}} \hat{V}(s, \mathbf{w}) \quad (11)$$

$$= \alpha (r + \gamma \hat{V}(s', \mathbf{w}) - \hat{V}(s, \mathbf{w})) \mathbf{x}(s_t) \quad (12)$$

Linear TD converges "close" to the global optimum (provable). This does not extend to non-linear TD (see Sutton & Barto, 2018).

From eval. to control: FA in GPI



1) Policy evaluation: Approximate policy evaluation $\hat{Q}(s, a, \mathbf{w}) \approx Q^\pi(s, a)$

2) Policy improvement: ϵ -greedy policy improvement

8) Deep-Q Learning

Use DL to replace hand-engineering of state space features w/ learning features from state data directly.

DQN (Mnih et al., 2015) Atari: Q(S,A) approx. by CNN w/ raw pixel inputs and discrete action outputs

Given our TD Q-learning update:

$$Q(s_t, a) \leftarrow Q(s_t, a) + \alpha [r_{t+1} + \gamma \max_{a'} Q(s_{t+1}, a') - Q(s_t, a)]$$

we want to learn $Q(s_{t+1}, a'; \mathbf{w})$ as a parametrised function (a neural network) with parameters \mathbf{w} .

We can define the TD error as our learning target that we want to reduce to zero.

$$\text{TD error}(w) = r_{t+1} + \gamma \max_{a'} Q(s_{t+1}, a'; \mathbf{w}) - Q(s_t, a; \mathbf{w})$$

Taking the gradient of E wrt \mathbf{w} we obtain:

$$\Delta \mathbf{w} = \alpha [r + \gamma \max_{a'} Q(s_{t+1}, a'; \mathbf{w}) - Q(s_t, a; \mathbf{w})] \nabla_{\mathbf{w}} Q(s_t, a; \mathbf{w})$$

Prblm w/ regular Deep-Q learning: maximisation bias = when taking max over all actions w/ (very) finite data, we may always overestimate values. **SOL:** Double Q-learning (DDQN)

Complication: The maximum Q value may be overestimated (variance-bias problem), because an unusually high value from the main network Q does not mean that there is an unusually high value from the target network Q' .

DDQN learning: more realistic (closer to final values), more stable, less biased (far less systematic overshooting).

Also see Double Dueling Q networks for Double-Q learning

Engineering needed to solve Atari games w/ DQN ("regular Deep-Q learning"):

1) **Experience Replay:** CNN overfitting latest experienced episodes. Inefficient use of interactive experience + highly correlated training samples (agent's recent actions generated from recent policy outputs).

SOL: experiences (traces) stored and replayed in minibatches to train on more than just last episode (makes data more iid); instead of running Q-learning on each s-a pair as they occur, experience replay buffer stores the traces sampled. Exp reuse (random sampled) => data-efficient (higher learning speed) + avoid catastrophic forgetting of R associated w/ replayed transitions

2) **Target Network:** unstable training (e.g. resonance effect, divergence) bc bootstrapping a continuous state space repr. **SOL:** slow learning down (resonance dampener) using target network Q' . Initialise main network Q and target network Q' . Use Q' to calculate TD-error. In frequently set $Q'=Q$ (params). Gives highly fluctuating Q time to settle (relaxation time) before updating Q'

3) **Clipping of Rewards:** all positive rewards = +1, all negative rewards = -1 to avoid diff reward scales making training unstable.

4) **Skipping of Frames:** reduce comp. cost and accelerate training time + make the game run at speed comparable to human RT by only using every 4 video game frames as input (60Hz => 15Hz).

DQN config (49 Atari 2600 games)

Input: 84x84x4 preprocessed image. 3 conv. => 2 FC (ReLU) => FC output layer RMSProp, minibatch=32, ϵ -greedy behav. $\pi(e=1->0.1)$ /Trained for 50million frames (38days game exp) + replay memory = 1million



DDQN (van Hasselt et al., 2017, AAAI): Use target network Q' for estimating best action selection; regular Q for estimating Q-value of s-a pair (or converse). Reduces frequency by which max Q-value may be overestimated bc less likely that both networks are overestimating the same action.

9) Policy Gradients. Policy-based methods: find π^* without V/Q fns. Faster convergence and often better for continuous and stochastic envs than value-based methods 1-8).

= look directly at parametrised π_{θ} w/ params θ . Optimise by looking at traces that a policy π_{θ} rolls out to correlate them w/ the rewards they incur.

Probability to observe a trace τ depends on (the policy weights θ) of π :

$$p(\tau) = p(s_1, a_1, \dots, s_T, a_T) = p(s_1) \pi_{\theta}(a_1 | s_1) p(s_2 | s_1, a_1) \pi_{\theta}(a_2 | s_2) p(s_3 | s_2, a_2) \dots \pi_{\theta}(a_T | s_T, a_{T-1})$$

Hard to measure τ Hard to model $\theta^* = \arg \max_{\theta} \mathbb{E} \left[\sum_{t=1}^T r(s_t, a_t) \right]_{\tau \sim p_{\theta}(\tau)}$

π^* obtained from θ^* instead of V^*/Q^* , where $\theta^* = \theta$ giving maximum average return.

Rationale: for continuous envs, infinite states/actions to estimate => value-based method = comp. expensive, esp. w/ GPI where policy improv step needs full scan of action space (max over A). A value fn may be used to learn θ but is not required for action selection.

Can use any parametric supervised ML model to learn $\pi(a|s)$; learned θ .

Use gradient ascent as we want to max performance.

Performance measure (episodic case)

$$J(\theta) = V^\pi(s_0) = \mathbb{E} \left[\sum_{t=1}^T r(s_t, a_t) \right]_{\tau \sim p_{\theta}(\tau)}$$

Approach 2: Direct Policy Gradients (using log trick)

1. average return approx. by empirical mean over N traces

$$J(\theta) = \mathbb{E} \left[\sum_{t=1}^T r(s_t, a_t) \right]_{\tau \sim p_{\theta}(\tau)} \approx \frac{1}{N} \sum_{i=1}^N \sum_t r(s_{i,t}, a_{i,t})$$

REINFORCE: insightful first-shot at PGs (popularising it)

1. sample $\{\tau^i\}$ from $\pi_{\theta}(a_t | s_t)$ (run the policy)

2. $\nabla_{\theta} J(\theta) \approx \sum_i \left(\sum_t r(s_{i,t}, a_{i,t}) \right) \left(\sum_t \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right)$

3. $\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$ In the limit of large amounts of data => model will converge to θ^*

No bias but high variance in the sampled trajectories => difficult to stabilise θ . Any erratic journey can cause suboptimal shift in the policy dist. Reduce variance (being smarter correlating rewards with trajectories) by subtracting a baseline from reward term (keep values smaller) or using an advantage term (Schuman et al., 2016) (see 10)).

Difficulty of computing policy gradient:

Depends on traces τ : meaning we need derivatives on action selection and the stationary dist of states $p(s)$, both determined by $\pi(a|s, \theta)$. Given env is also generally unknown, difficult to estimate the effect of a policy update on state dist.

2. log trick to derive gradient

$$\nabla_{\theta} J(\theta) = \int \nabla_{\theta} \pi_{\theta}(\tau) r(\tau) d\tau = \int \pi_{\theta}(\tau) \nabla_{\theta} \log \pi_{\theta}(\tau) r(\tau) d\tau = \mathbb{E}_{\tau \sim \pi_{\theta}(\tau)} [\nabla_{\theta} \log \pi_{\theta}(\tau) r(\tau)]$$

$$\nabla_{\theta} \log \pi_{\theta}(\tau) = \nabla_{\theta} \left[\log p(s_1) + \sum_t \log \pi_{\theta}(a_t | s_t) + \log p(s_{t+1} | s_t, a_t) \right]$$

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}(\tau)} \left[\left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right) \left(\sum_{t=1}^T r(s_t, a_t) \right) \right]$$

PG = trial-and-error like MC learning

= one version of the **Policy Gradient Theorem** (policy-centered equivalent of Bellman Theorem. Implemented in REINFORCE

10) Actor-Critic Methods => model-learning to improve the model (!= direct RL improving V/Q/ π , as in 1-9)). A-C = improving V/Q/ π via model = indirect/model-based RL => get better π w/ fewer interactions. Can have both indirect and direct methods => Split RL model into an actor (compute a based on s) and critic (produce Q-values of s, a = "model"). Actor: input=s, output=a. Controls how agent behaves by learning π^* (policy-based learning). Critic: evaluates a (value-based learning). The two models "compete" and each gets better at its role: key point = combined arch learns better than the 2 separate networks would individually.

A simple Q-driven policy-gradient actor-critic.

Algorithm 1 Q Actor-Critic

Initialize parameters s, θ, \mathbf{w} and learning rates α, β ; sample $a \sim \pi_{\theta}(a|s)$.
for $t = 1 \dots T$: do
Sample reward $r_t \sim R(s_t, a)$ and next state $s' \sim P(s'|s, a)$
Then sample the next action $a' \sim \pi_{\theta}(a'|s')$
Update the policy parameters: $\theta \leftarrow \theta + \alpha \nabla_{\theta} Q_{\mathbf{w}}(s, a) \nabla_{\theta} \log \pi_{\theta}(a|s)$; Compute the correction (TD error) for action-value at time t:
 $\delta_t = r_t + \gamma Q_{\mathbf{w}}(s', a') - Q_{\mathbf{w}}(s, a)$
and use it to update the parameters of Q function:
 $\mathbf{w} \leftarrow \mathbf{w} + \alpha \delta_t \nabla_{\mathbf{w}} Q_{\mathbf{w}}(s, a)$
Move to a $\leftarrow a'$ and $s \leftarrow s'$
end for

Advantage A-C (A2C): learn A, not Q: $Q(s, a) = V(s) + A(s, a)$

Eval. of action a based on how much it improves s value. A2C reduces high var. of actor (=PG method) w/o adding bias, stabilises model in training

Asynchronous (A3C): multiple indpt agents (networks) interact w/ diff env copy in // => explore bigger part of S-A space in much less time. Trained in // and update periodically (asynchronously) a global network holding shared params. After each update, each agent copies global params + resume indpt exploring until next update (cool if large scale simulators). Info flow btw agents and global network, and btw agents (given param reset).

DDPG: deterministic policy, model-free, off-policy, A-C. Continuous actions version of DQN. Uses bootstrapping to learn Q fn, and to learn π from estimated Q fn. Explore in continuous space using Gaussian policy (~= discrete action ϵ -greediness). Replay buffer, minibatch, target network (=DQN) control learning variability. MSBE. PG-based actor. Smooth updates to A&C target networks $\theta^{t+1} = \beta \theta^t + (1 - \beta) \theta^t, 0 < \beta \ll 1$